

# Toward Visually Robust Quadruped Parkour via Monocular Depth Estimation

Theo Lemay  
*Undergraduate Thesis*  
theo.lemay@queensu.ca

April 2026

## Abstract

Vision-based quadruped parkour typically relies on onboard depth cameras, limiting deployment to robots with specialized sensors. We investigate whether a pre-trained *monocular depth estimator* can replace depth cameras for locomotion by producing a canonical depth representation from RGB, without photorealistic rendering or domain randomization.

We present a two-stage pipeline: a privileged teacher trained via PPO with ground-truth terrain scans, distilled into a visual student via DAgger operating on depth images from Depth Anything V2 (DA2). In simulation, this pipeline enables basic parkour on stairs, hurdles, and gaps under conditions characterized by a *Depth Signal Quality* (DSQ) metric. Visual robustness experiments show that DA2-based students tolerate moderate lighting and color perturbations within simulation, though performance degrades on featureless surfaces and under combined perturbations. All results are in simulation; real-world transfer remains future work.

## 1 Introduction

Legged robots that traverse irregular terrain—stairs, hurdles, and gaps—must perceive geometry and coordinate vision and control at high rates. Training locomotion policies in simulation with reinforcement learning (RL) and transferring them to real hardware has proven effective for agile skills [8, 11, 15]. Recent work [2, 4, 7, 24] has extended this approach to parkour—climbing, jumping, and crawling—through teacher-student distillation. In these systems, a *teacher* trains with privileged terrain information (e.g. ground-truth heightmaps), and a *student* reproduces the teacher’s behavior from onboard sensors.

A central assumption, however, is access to an onboard depth sensor. Depth-camera policies transfer well from simulation on platforms that provide depth sensing [1, 12], but not all robots do. The Unitree Go2, for instance, ships with RGB cameras and a LiDAR sensor, but no onboard depth camera [20].

Deploying vision policies from RGB therefore requires bridging the *sim-to-real gap*. Three strategies have emerged: (1) **domain randomization** [19], which randomizes visual properties; (2) **generative augmentation** [23], which uses diffusion models for photorealistic training images; and (3) **canonical representations**, which project both domains into a shared space via a foundation model.

We pursue strategy (3) using Depth Anything V2 (DA2) [22], a pre-trained monocular depth estimator trained on millions of diverse real images. Our premise is that if DA2 maps RGB into a depth space that preserves terrain geometry across visual conditions, then a student policy operating on estimated depth may inherit partial robustness to appearance variation without photorealistic rendering or explicit domain randomization.

**Thesis statement.** We hypothesize that pre-trained monocular depth estimation can provide a useful intermediate representation for vision-based quadruped parkour, enabling a student trained on estimated depth in simulation to tolerate a range of visual conditions without explicit domain randomization. We evaluate this hypothesis entirely in simulation; real-world transfer remains future work.

### Contributions.

- A two-stage teacher-student pipeline from privileged scandot heightmaps to monocular depth for quadruped parkour.
- A demonstration that estimated depth enables basic parkour in simulation.
- Identification of depth signal quality—governed by model scale and terrain texture—as the key factor limiting performance.
- A visual robustness study showing that DA2-based students tolerate lighting and color shifts but depend on terrain texture for depth signal quality.

## 2 Related Work

**RL for legged locomotion.** Modern quadruped controllers are typically learned via reinforcement learning

in simulation and transferred to real hardware [8, 15]. Chen et al. [3] introduced the privileged teacher–student paradigm, where a teacher with access to ground-truth state is distilled into a sensorimotor student. Lee et al. [11] applied this to locomotion, training a student from proprioceptive history alone and achieving zero-shot sim-to-real transfer on rough terrain. Kumar et al. [10] introduced rapid motor adaptation, using a learned adaptation module to estimate privileged parameters online. Miki et al. [12] extended this framework to perceptive locomotion, incorporating exteroceptive depth inputs.

**Vision-based parkour.** Agarwal et al. [1] introduced the scandot representation for egocentric depth-based locomotion, distilling a privileged teacher into a visual student. Extreme Parkour [4] and Robot Parkour Learning [24] build on this with CNN–GRU students that process depth images, demonstrating climbing, jumping, and crawling on real quadrupeds. SoloParkour [2] replaces distillation with off-policy RL warm-started from privileged experience, and ANYmal Parkour [7] uses a hierarchical skill-selection policy. All five systems assume access to a depth camera at deployment.

Jenkins [9] is the closest work to ours, training an RGB parkour policy using visual domain randomization with a fine-tuned MobileNetV2 backbone. In real-world trials, Jenkins succeeded on hurdles but failed on gaps, suggesting that domain randomization alone may not fully close the visual sim-to-real gap. We pursue a complementary strategy: instead of learning invariance through randomized exposure, we project RGB into a canonical depth space using a pre-trained estimator.

**Sim-to-real for vision.** Domain randomization [19] perturbs visual properties (textures, lighting, colors) so the policy learns invariance through exposure. LucidSim [23] replaces randomization with generative models that produce photorealistic training images. Xue et al. [21] combine photorealistic simulation with a teacher–student–bootstrap framework for humanoid locomanipulation from RGB. All three approaches require substantial rendering or generation overhead. We instead project RGB into a canonical depth space using a pre-trained estimator.

**Monocular depth estimation.** MiDaS [13] pioneered monocular depth estimation across diverse datasets. Depth Anything V2 (DA2) [22] achieves state-of-the-art zero-shot depth estimation using a vision transformer encoder trained on 142M images. DA2 offers both relative and metric variants; we use the metric model, which predicts absolute depth in meters. Figure 1 illustrates the qualitative behavior of DA2 across diverse scenes and motivates its use here as a canonical visual representation.



Figure 1: Depth Anything V2 [5, 22].

## 3 Mathematical Foundations

### 3.1 Problem Formulation

We model the parkour task as a Partially Observable Markov Decision Process  $(\mathcal{S}, \mathcal{A}, \mathcal{O}, T, R, \gamma)$  with state space  $\mathcal{S}$ , action space  $\mathcal{A}$ , observation space  $\mathcal{O}$ , transition dynamics  $T$ , reward function  $R$ , and discount factor  $\gamma$ . The agent receives partial observation  $o_t \in \mathcal{O}$  and selects action  $a_t \in \mathcal{A} = \mathbb{R}^{12}$  (target joint position offsets). Observations decompose into three groups:

- **Proprioception**  $o_{\text{prop}} \in \mathbb{R}^{53}$ : base angular velocity, projected gravity, velocity and heading commands, joint positions, joint velocities, previous actions, and foot contacts.
- **Exteroception**  $o_{\text{ext}}$ : scandot heightmaps ( $\mathbb{R}^{132}$ ) for the teacher, or depth images for the student. Scandots [1] are terrain height samples on a  $12 \times 11$  grid in the robot’s body frame. Each value records the height difference between the base and the terrain surface, clipped to  $[-1, 1]$ .
- **Privileged**  $o_{\text{priv}}$ : friction and other physics parameters available only during teacher training.

### 3.2 PPO

The teacher policy  $\pi_\theta$  is trained with Proximal Policy Optimization (PPO) [17], which maximizes a clipped surrogate objective:

$$L^{\text{CLIP}} = \mathbb{E}_t \left[ \min \left( r_t \hat{A}_t, \text{clip}(r_t, 1-\epsilon, 1+\epsilon) \hat{A}_t \right) \right], \quad (1)$$

where  $r_t = \pi_\theta(a_t|o_t)/\pi_{\theta_{\text{old}}}(a_t|o_t)$  is the probability ratio and  $\hat{A}_t$  is the GAE [16] advantage estimate.

### 3.3 Reward Function

The teacher uses a shaped reward that combines task progress with regularization for stable locomotion. Positive terms reward tracking the commanded velocity and heading, making forward progress, and reaching waypoints and course completion. Penalty terms discourage early termination, collisions, unstable body motion, large torques and joint accelerations, abrupt action changes, and joint or torque limit violations. See the full reward function in Table 4.




Figure 2: Two-stage training pipeline. The teacher uses privileged scandot observations; the student replaces these with either GT depth or DA2-inferred depth.

### 3.4 DAgger Distillation

The student policy is trained via Dataset Aggregation (DAgger) [14]. At each step the student observes  $o_{\text{prop}}$  and a depth image and predicts  $a_{\text{student}}$ , but the teacher’s action  $a_{\text{teacher}}$  is executed in the environment. The student minimizes:

$$\mathcal{L} = \underbrace{\|a_{\text{teacher}} - a_{\text{student}}\|_2}_{\text{action loss}} + \underbrace{\|y_{\text{oracle}} - \hat{y}\|_2}_{\text{yaw loss}}, \quad (2)$$

where  $\hat{y} \in \mathbb{R}^2$  is the student’s predicted heading and  $y_{\text{oracle}}$  is the ground-truth heading.

## 4 System Architecture

### 4.1 Overview

All experiments use the Genesis GPU simulator [6] with the legged\_gym framework [15] and rsl\_rl [18] on 4× NVIDIA L40S GPUs. The robot is a Unitree Go2 quadruped with 12 actuated joints. Actions  $a_t \in \mathbb{R}^{12}$  are target joint positions, updated at 50 Hz and converted to torques by a PD controller  $\tau = K_p(a_t - q) - K_d\dot{q}$ .

The system follows the two-stage privileged distillation pipeline described in Section 1: a teacher trains with ground-truth scandot heightmaps [1], then a student reproduces the teacher’s behavior from GT depth or DA2-estimated depth via DAgger (Figure 2).

### 4.2 Teacher

The teacher policy  $\pi_T$  maps proprioception, privileged physics parameters, and scandot heightmaps to joint targets. Three parallel MLPs encode the observation groups into 32-dim latent vectors, which are concatenated with proprioception and fed through the actor MLP  $\mathbb{R}^n \rightarrow \mathbb{R}^{12}$  to produce joint position offsets.




Figure 3: Student visual encoder. The depth CNN and proprioception are fused and processed by a GRU that produces the terrain latent and heading estimate used by the teacher’s actor MLP.




Figure 4: Parkour curriculum progression across obstacle families. Difficulty increases from left to right by increasing step count/height, hurdle size, and gap width.

### 4.3 Student

The student replaces the scandot encoder with a visual encoder that recovers the terrain latent  $z \in \mathbb{R}^{32}$  from a depth image (Figure 3). A CNN maps the depth image  $d \in \mathbb{R}^{58 \times 87}$  to a 32-dim feature, which is concatenated with proprioception and compressed by an MLP. A GRU integrates these fused features across timesteps, producing  $z$  and a heading estimate  $\hat{y} \in \mathbb{R}^2$ . The teacher’s actor MLP is reused, taking  $[o_{\text{prop}}; z]$  to output joint targets. During DAgger distillation, the student minimizes the loss in Eq. 2 while the teacher’s actions are executed in the environment.

### 4.4 Terrain and Curriculum

The simulator generates parkour terrains with three obstacle families (stairs, hurdles, gaps) at four difficulty levels per family (Figure 4) with gap widths from 0.24 m to 0.50 m. These obstacles are simpler than those in prior work [2, 4, 24]; our goal is to establish a baseline for evaluating the DA2 depth pipeline rather than to push the limits of agile locomotion. During teacher training, a curriculum advances robots to harder rows upon reaching a success threshold.

Table 1: Baseline success rates on gap terrain (%).

|                     | Row 0 | Row 1 | Row 2 | Row 3 |
|---------------------|-------|-------|-------|-------|
| Scandots teacher    | 98.5  | 100.0 | 99.0  | 98.5  |
| GT-depth student    | 94.5  | 94.5  | 89.1  | 87.5  |
| Zero-depth ablation | 8.0   | 2.0   | 0.0   | 0.0   |

## 5 Experiments and Results

We evaluate the DA2 pipeline through three experiments: (1) establishing upper-bound baselines with ground-truth sensing, (2) closing the gap between estimated and ground-truth depth by varying model capacity and terrain texture, and (3) testing visual robustness under visual perturbations. We report *success rate* (fraction of episodes completing the full course) and *progress ratio* (mean fractional distance along the course before termination), both averaged over 256 episodes.

For visual robustness, we also report *depth signal quality* (DSQ), defined as the scale- and offset-invariant Pearson correlation between ground-truth and DA2-estimated depth images, averaged over environments and time steps.

$$\text{DSQ} = \frac{1}{T} \sum_{t=1}^T \frac{1}{N} \sum_{i=1}^N \frac{\sum_p (d_{i,p}^{\text{GT}} - \bar{d}_i^{\text{GT}})(d_{i,p}^{\text{DA2}} - \bar{d}_i^{\text{DA2}})}{\|d_i^{\text{GT}} - \bar{d}_i^{\text{GT}}\|_2 \|d_i^{\text{DA2}} - \bar{d}_i^{\text{DA2}}\|_2} \quad (3)$$

### 5.1 Baselines

We first establish upper bounds for the privileged teacher and the ground-truth depth sensing student (GT-depth). The privileged teacher achieves near-perfect success across all gap difficulties, shown in Table 1.

The GT-depth student, receiving simulator depth images, achieves success on all rows, confirming the CNN-GRU architecture extracts sufficient geometry from a single depth viewpoint. A zero-depth ablation (all depth pixels set to zero) drops success to 8%, verifying the policy relies on depth rather than exploiting proprioceptive heuristics.

### 5.2 Parkour from Estimated Depth

We next test whether parkour is possible from estimated depth. Starting from the GT-depth setup, we replace simulator depth with DA2-small. As shown in Table 2, DA2-small succeeds on the early rows but fails completely on Rows 2 and 3. Switching to DA2-base reaches Row 3, but still falls well short of GT-depth on the hardest rows.

The remaining limitation is the visual input. On untextured terrain, the gap boundary provides weak cues, so

the depth estimate is ambiguous. As shown in Table 2, adding a checkerboard terrain texture makes the gap geometry visible and closes most of the remaining gap to GT-depth. Figure 6 shows examples of depth estimation across representative surface conditions.

Table 2: Policy success rate (%) as model scale and terrain texture vary.

| Configuration      | Row 0 | Row 1 | Row 2 | Row 3 |
|--------------------|-------|-------|-------|-------|
| DA2-small          | 56.3  | 33.6  | —     | —     |
| DA2-base           | 68.4  | 83.2  | 26.6  | 0.8   |
| DA2-base + texture | 87.5  | 93.0  | 87.1  | 80.5  |
| GT-depth baseline  | 94.5  | 94.5  | 89.1  | 87.5  |

Figure 5 shows the training dynamics. Reward rises steeply as the policies master flat terrain, then plateaus as the curriculum exposes harder obstacles. The textured model surpasses the untextured models and approaches the GT-depth baseline.




Figure 5: Success rate during student training. Breaking success down by obstacle type shows that most of the gain comes from gap traversal; stairs and hurdles are less sensitive to texture because their geometry remains visible.

### 5.3 Visual Robustness

We next test robustness to the appearance shifts that would arise in sim-to-real transfer by perturbing the trained student only at *test time*. All perturbations are applied within simulation; these results therefore measure within-sim robustness, not actual sim-to-real transfer. Table 3 reports success across six terrain configurations (three obstacle families  $\times$  two difficulty levels) for texture replacement, lighting/color changes, and combined perturbations, while Figure 7 plots DSQ against policy success.




Figure 6: Robot-view comparison across representative terrain textures. Each panel shows rendered RGB (left) and the corresponding DA2 depth estimate (right). Featureless terrain produces the weakest obstacle localization and the lowest DSQ, while checkerboard, grass, and wood all preserve substantially more geometric structure.




Figure 7: Success rate vs. DSQ across all six terrain configurations. Color denotes obstacle type and difficulty; marker shape denotes perturbation category. Featureless surfaces cluster at low DSQ with catastrophic failure; most textured conditions maintain  $DSQ > 0.85$  but success varies by difficulty.

The results reveal a clear obstacle-geometry hierarchy. At the easier setting, stairs and hurdles are largely insensitive to appearance changes because strong geometric edges remain visible across textures and lighting. Gaps are more sensitive because their boundary is visible mainly through surface texture, so performance depends more strongly on the quality of the depth estimate.

At the harder setting, all three obstacle families degrade, including stairs and hurdles. Featureless surfaces cause the largest failures, while lighting and color shifts have a smaller effect. This pattern suggests DA2 is robust to photometric variation but still depends on texture when obstacle boundaries are weak.

As Figure 7 shows, DSQ is a useful failure indicator: featureless surfaces collapse both DSQ and success. However, high DSQ alone does not guarantee robustness, since some unseen textures still reduce success despite preserving coarse depth structure.

Table 3: Visual robustness: test-time success (%) under visual perturbations across obstacle types.

| Perturbation                      | Gaps |      | Stairs |      | Hurdles |      |
|-----------------------------------|------|------|--------|------|---------|------|
|                                   | R0   | R3   | R0     | R3   | R0      | R3   |
| Baseline                          | 96.1 | 94.1 | 100    | 92.2 | 96.9    | 94.9 |
| <i>Realistic surface textures</i> |      |      |        |      |         |      |
| Concrete                          | 79.3 | 12.1 | 99.6   | 25.8 | 95.7    | 73.0 |
| Stone tiles                       | 93.0 | 78.9 | 100    | 96.9 | 95.3    | 90.6 |
| Gravel                            | 91.4 | 48.8 | 100    | 93.4 | 93.4    | 34.8 |
| <i>Structured / featureless</i>   |      |      |        |      |         |      |
| Noise                             | 82.8 | 29.3 | 100    | 95.3 | 97.3    | 47.7 |
| No texture                        | 16.8 | 0.0  | 100    | 21.9 | 91.4    | 18.2 |
| <i>Brightness / color</i>         |      |      |        |      |         |      |
| Dim (0.3 $\times$ )               | 82.8 | 58.2 | 99.6   | 94.1 | 99.6    | 95.7 |
| Color $\pm 30\%$                  | 79.3 | 32.0 | 99.6   | 94.9 | 97.3    | 87.9 |
| <i>Combined*</i>                  |      |      |        |      |         |      |
| Mild                              | 76.6 | 8.2  | 100    | 15.2 | 97.3    | 79.3 |
| Extreme                           | 34.0 | 0.4  | 99.6   | 66.8 | 89.5    | 19.1 |

\* Mild: concrete + dim (0.6 $\times$ ) + brightness jitter  $\pm 20\%$ . Extreme: noise + dim (0.4 $\times$ ) + low contrast ( $\gamma=1.8$ ) + color jitter  $\pm 20\%$ .

## 6 Discussion and Conclusion

### 6.1 Summary

We presented a teacher-student pipeline for vision-based quadruped parkour that uses monocular depth estimation as an intermediate representation. The GT-depth baseline shows that the student architecture can solve the task when accurate geometry is available. The DA2 experiments show that the main bottleneck is the quality of the recovered depth signal rather than the policy architecture. Increasing model capacity and adding structured terrain texture closes most of the gap to the GT-depth baseline on the gap curriculum.

The visual robustness experiments further refine this result. The DA2-based student tolerates moderate lighting and color changes and transfers to several unseen textured surfaces without photorealistic rendering or explicit domain randomization. But the method is not



appearance-invariant in general: performance degrades on harder settings and collapses on featureless surfaces, where monocular depth cannot recover reliable geometry. DSQ is a useful indicator of whether the visual pipeline preserves terrain structure, but it does not fully determine policy success. Overall, the results support monocular depth as a plausible bridge from RGB to locomotion, but actual sim-to-real transfer remains unvalidated.

## 6.2 Limitations

Several limitations constrain this conclusion. First, all experiments are in simulation, so the central sim-to-real claim is still unverified on hardware. Second, the task scope is modest: the obstacles are simpler than those in recent real-world parkour systems, and the visual robustness study covers two difficulty levels across three families rather than the full curriculum. Third, the method depends on a single off-the-shelf depth estimator. DA2 degrades on weak-texture terrain, and high DSQ does not explain all failures on unseen textures. Fourth, DA2 processes each frame independently, so depth estimates may flicker between consecutive timesteps. The GRU in the student encoder may partially smooth such temporal noise, but we did not measure frame-to-frame depth consistency or test temporally consistent depth models, leaving this an open question for deployment at high control rates. Fifth, our visual robustness study covers texture, lighting, and color perturbations, but real-world deployment introduces additional failure modes that monocular depth estimators are known to struggle with: transparent or reflective surfaces, thin structures such as railings, motion blur during dynamic gaits, and close-range objects at the limits of the model’s training distribution. These are likely to arise in outdoor parkour environments and may compound the texture sensitivity already observed. Finally, the perception stack is still heavy for onboard deployment. The current system is therefore best viewed as a proof of concept rather than a deployable parkour pipeline.

## 6.3 Future Work

The next step is real-world evaluation on the Unitree Go2. This will test robustness under real camera noise, motion blur, exposure shifts, and outdoor lighting, and determine whether DSQ remains useful as a deployment diagnostic. A second direction is broader evaluation across stairs, hurdles, and more difficult compound courses. On the modeling side, smaller or temporally consistent depth estimators could reduce compute cost and stabilize predictions, while joint adaptation of the depth estimator and student policy may improve robustness on harder textures and further close the gap to the GT-depth baseline.

## References

- [1] Ananye Agarwal, Ashish Kumar, Jitendra Malik, and Deepak Pathak. Legged locomotion in challenging terrains using egocentric vision. In *Conference on robot learning*, pages 403–415. PMLR, 2023.
- [2] Elliot Chane-Sane, Joseph Amigo, Thomas Flayols, Ludovic Righetti, and Nicolas Mansard. Soloparkour: Constrained reinforcement learning for visual locomotion from privileged experience. In *Proceedings of The 8th Conference on Robot Learning*, volume 270 of *Proceedings of Machine Learning Research*, pages 4268–4285. PMLR, 2025.
- [3] Dian Chen, Brady Zhou, Vladlen Koltun, and Philipp Krähenbühl. Learning by cheating. In *Conference on Robot Learning (CoRL)*, pages 66–75. PMLR, 2020.
- [4] Xuxin Cheng, Kexin Shi, Ananye Agarwal, and Deepak Pathak. Extreme parkour with legged robots. In *IEEE International Conference on Robotics and Automation (ICRA)*, 2024.
- [5] Depth Anything Authors. Depth anything V2 official repository, 2024. URL <https://github.com/DepthAnything/Depth-Anything-V2>.
- [6] Genesis Authors. Genesis: A generative and universal physics engine for robotics and beyond, 2024. URL <https://github.com/Genesis-Embodied-AI/Genesis>.
- [7] David Hoeller, Nikita Rudin, Dhionis Sako, and Marco Hutter. Anymal parkour: Learning agile navigation for quadrupedal robots. *Science Robotics*, 2024.
- [8] Jemin Hwangbo, Joonho Lee, Alexey Dosovitskiy, Dario Bellicoso, Vassilios Tsounis, Vladlen Koltun, and Marco Hutter. Learning agile and dynamic motor skills for legged robots. *Science Robotics*, 4(26):eaau5872, 2019.
- [9] Andrew Jenkins. Learning sim-to-real robot parkour from RGB images. M.eng. thesis, Massachusetts Institute of Technology, 2024. URL <https://dspace.mit.edu/handle/1721.1/156972>.
- [10] Ashish Kumar, Zipeng Fu, Deepak Pathak, and Jitendra Malik. RMA: Rapid motor adaptation for legged robots. In *Robotics: Science and Systems (RSS)*, 2021.
- [11] Joonho Lee, Jemin Hwangbo, Lorenz Wellhausen, Vladlen Koltun, and Marco Hutter. Learning quadrupedal locomotion over challenging terrain. *Science robotics*, 5(47):eabc5986, 2020.

- [12] Takahiro Miki, Joonho Lee, Jemin Hwangbo, Lorenz Wellhausen, Vladlen Koltun, and Marco Hutter. Learning robust perceptive locomotion for quadrupedal robots in the wild. *Science Robotics*, 7(62):eabk2822, 2022.
- [13] René Ranftl, Katrin Lasinger, David Hafner, Konrad Schindler, and Vladlen Koltun. Towards robust monocular depth estimation: Mixing datasets for zero-shot cross-dataset transfer. *IEEE Transactions on Pattern Analysis and Machine Intelligence*, 44(3):1623–1637, 2022. doi: 10.1109/TPAMI.2020.3019967.
- [14] Stéphane Ross, Geoffrey J. Gordon, and Drew Bagnell. A reduction of imitation learning and structured prediction to no-regret online learning. In *Proceedings of the Fourteenth International Conference on Artificial Intelligence and Statistics*, pages 627–635. JMLR Workshop and Conference Proceedings, 2011.
- [15] Nikita Rudin, David Hoeller, Philipp Reist, and Marco Hutter. Learning to walk in minutes using massively parallel deep reinforcement learning. In *Conference on Robot Learning (CoRL)*, pages 91–100. PMLR, 2022.
- [16] John Schulman, Philipp Moritz, Sergey Levine, Michael I. Jordan, and Pieter Abbeel. High-dimensional continuous control using generalized advantage estimation. In *International Conference on Learning Representations (ICLR)*, 2016.
- [17] John Schulman, Filip Wolski, Prafulla Dhariwal, Alec Radford, and Oleg Klimov. Proximal policy optimization algorithms, 2017.
- [18] Clemens Schwarke, Mayank Mittal, Nikita Rudin, David Hoeller, and Marco Hutter. RSL-RL: A learning library for robotics research. *arXiv preprint arXiv:2509.10771*, 2025.
- [19] Josh Tobin, Rachel Fong, Alex Ray, Jonas Schneider, Wojciech Zaremba, and Pieter Abbeel. Domain randomization for transferring deep neural networks from simulation to the real world. In *IEEE/RSJ International Conference on Intelligent Robots and Systems (IROS)*, 2017.
- [20] Unitree Robotics. Go2 — Unitree quadruped robot, 2024. URL <https://www.unitree.com/go2>. Accessed April 2026.
- [21] Haoru Xue, Tairan He, Zi Wang, Qingwei Ben, Wenli Xiao, Zhengyi Luo, Xingye Da, Fernando Castañeda, Guanya Shi, Shankar Sastry, Linxi Fan, and Yuke Zhu. Opening the sim-to-real door for humanoid pixel-to-action policy transfer, 2025.
- [22] Lihe Yang, Bingyi Kang, Zilong Huang, Zhen Zhao, Xiaogang Xu, Jiashi Feng, and Hengshuang Zhao. Depth anything V2, 2024.
- [23] Alan Yu, Ge Yang, Ran Choi, Yajvan Ravan, John J. Leonard, and Phillip Isola. Learning visual parkour from generated images. In *Proceedings of The 8th Conference on Robot Learning*, volume 270 of *Proceedings of Machine Learning Research*, pages 2500–2516. PMLR, 2025.
- [24] Ziwen Zhuang, Zipeng Fu, Jianren Wang, Christopher Atkeson, Sören Schwertfeger, Chelsea Finn, and Hang Zhao. Robot parkour learning. In *Conference on Robot Learning (CoRL)*, 2023.

## A Hyperparameters

This appendix lists the training hyperparameters used in the experiments. Unless noted otherwise, values are shared across all runs.

### A.1 Reward Function

Table 4 lists the reward terms and scale factors used for teacher training. The reward at each step is  $r_t = \sum_i w_i r_i(s_t, a_t)$ , where  $w_i$  is the scale and  $r_i$  is the per-term reward signal.

Table 4: Teacher reward terms and scales.

| Term                               | Scale                 |
|------------------------------------|-----------------------|
| <i>Task progress</i>               |                       |
| Tracking goal velocity             | 3.0                   |
| Tracking goal heading              | 1.5                   |
| Progress along course              | 2.0                   |
| Waypoint reached bonus             | 20.0                  |
| Success bonus                      | 40.0                  |
| <i>Termination &amp; collision</i> |                       |
| Termination                        | −50.0                 |
| Collision                          | −2.0                  |
| Feet stumble                       | −1.0                  |
| Feet edge                          | −2.0                  |
| <i>Locomotion regularization</i>   |                       |
| Torques                            | $-2.0 \times 10^{-4}$ |
| Joint power                        | $-2.0 \times 10^{-4}$ |
| Joint acceleration                 | $-2.0 \times 10^{-7}$ |
| Action rate                        | −0.01                 |
| Action smoothness                  | −0.01                 |
| Joint position limits              | −2.0                  |
| Torque limits                      | −0.2                  |
| <i>Body stability</i>              |                       |
| Linear velocity $z$                | −0.5                  |
| Angular velocity $xy$              | −0.05                 |
| Orientation                        | −0.5                  |
| Flat back                          | −3.0                  |

## A.2 PPO (Teacher)

Table 5: PPO hyperparameters for teacher training.

| Parameter                  | Value                |
|----------------------------|----------------------|
| Environments               | 4096                 |
| Steps per env per update   | 48                   |
| Mini-batches               | 4                    |
| Learning epochs per update | 5                    |
| Learning rate              | $1 \times 10^{-3}$   |
| LR schedule                | adaptive (target KL) |
| Desired KL                 | 0.01                 |
| Discount $\gamma$          | 0.99                 |
| GAE $\lambda$              | 0.95                 |
| Clip $\epsilon$            | 0.2                  |
| Entropy coefficient        | 0.01                 |
| Value loss coefficient     | 1.0                  |
| Max gradient norm          | 1.0                  |
| Max iterations             | 2500                 |

Table 7: Network architecture details.

| Component                    | Specification                                                                                                                    |
|------------------------------|----------------------------------------------------------------------------------------------------------------------------------|
| <i>Teacher</i>               |                                                                                                                                  |
| Scandot encoder              | [128, 64, 32] MLP                                                                                                                |
| Actor MLP                    | [512, 256, 128] $\rightarrow$ 12                                                                                                 |
| Critic MLP                   | [1024, 512, 256] $\rightarrow$ 1                                                                                                 |
| Activation                   | ELU                                                                                                                              |
| <i>Student depth encoder</i> |                                                                                                                                  |
| CNN                          | Conv2d(1,32,5) $\rightarrow$ MaxPool(2) $\rightarrow$<br>Conv2d(32,64,3) $\rightarrow$ FC(62400,128)<br>$\rightarrow$ FC(128,32) |
| Input resolution             | $58 \times 87$ (H $\times$ W)                                                                                                    |
| Combination MLP              | $[d_{\text{cnn}} + d_{\text{prop}}, 128, 32]$                                                                                    |
| GRU                          | hidden dim = 512, 1 layer                                                                                                        |
| Output                       | [512, 34] $\rightarrow$ Tanh<br>(32 latent + 2 yaw)                                                                              |
| Actor MLP (reused)           | [512, 256, 128] $\rightarrow$ 12                                                                                                 |

## A.3 DAgger Distillation (Student)

Table 6: DAgger hyperparameters for student distillation.

| Parameter                    | Value              |
|------------------------------|--------------------|
| Environments                 | 48                 |
| Steps per env per update     | 120                |
| BPTT window                  | 24                 |
| Learning rate                | $2 \times 10^{-3}$ |
| Actor LR scale (DA2 student) | 0.1                |
| LR schedule (DA2 student)    | cosine             |
| Action loss coefficient      | 1.0                |
| Yaw loss coefficient         | 1.0                |
| Yaw threshold (MTS)          | 0.6                |
| Max gradient norm            | 1.0                |
| Max iterations               | 5000               |

Table 8: PD control and domain randomization parameters.

| Parameter                   | Value            |
|-----------------------------|------------------|
| <i>PD controller</i>        |                  |
| $K_p$                       | 30.0             |
| $K_d$                       | 0.75             |
| Action scale                | 0.25             |
| Control dt                  | 0.02 s (50 Hz)   |
| Decimation                  | 4                |
| <i>Domain randomization</i> |                  |
| Friction range              | [0.2, 1.7]       |
| Added base mass             | $[-1.0, 1.0]$ kg |
| CoM displacement $xyz$      | $\pm 0.03$ m     |
| $K_p$ / $K_d$ scale         | [0.8, 1.2]       |
| Push interval               | 15 s             |
| Max push velocity           | 1.0 m/s          |

## A.4 Network Architecture

## A.5 Control and Domain Randomization

## A.6 Terrain and Curriculum

Terrain sections are 12.0 m long and 3.6 m wide. The curriculum grid has 4 rows (difficulty levels)  $\times$  12 columns (3 families  $\times$  4 variants). Robots advance to harder rows when their success rate exceeds 85% and regress when it falls below 40%.

Table 9: Obstacle dimensions per curriculum row.

|                 | Row 0 | Row 1 | Row 2 | Row 3 |
|-----------------|-------|-------|-------|-------|
| <i>Stairs</i>   |       |       |       |       |
| Step height (m) | 0.06  | 0.08  | 0.10  | 0.12  |
| Step count      | 2     | 3     | 4     | 4     |
| <i>Hurdles</i>  |       |       |       |       |
| Height (m)      | 0.10  | 0.14  | 0.20  | 0.26  |
| Length (m)      | 0.20  | 0.22  | 0.25  | 0.28  |
| <i>Gaps</i>     |       |       |       |       |
| Width (m)       | 0.24  | 0.32  | 0.42  | 0.50  |